

cs231n Lecture 8

Deep Learning Software

2176410 한현경

Contents

01. CPU vs. GPU

02. Deep Learning Frameworks

(1) Tensorflow

(2) PyTorch

(3) Caffe/Caffe2

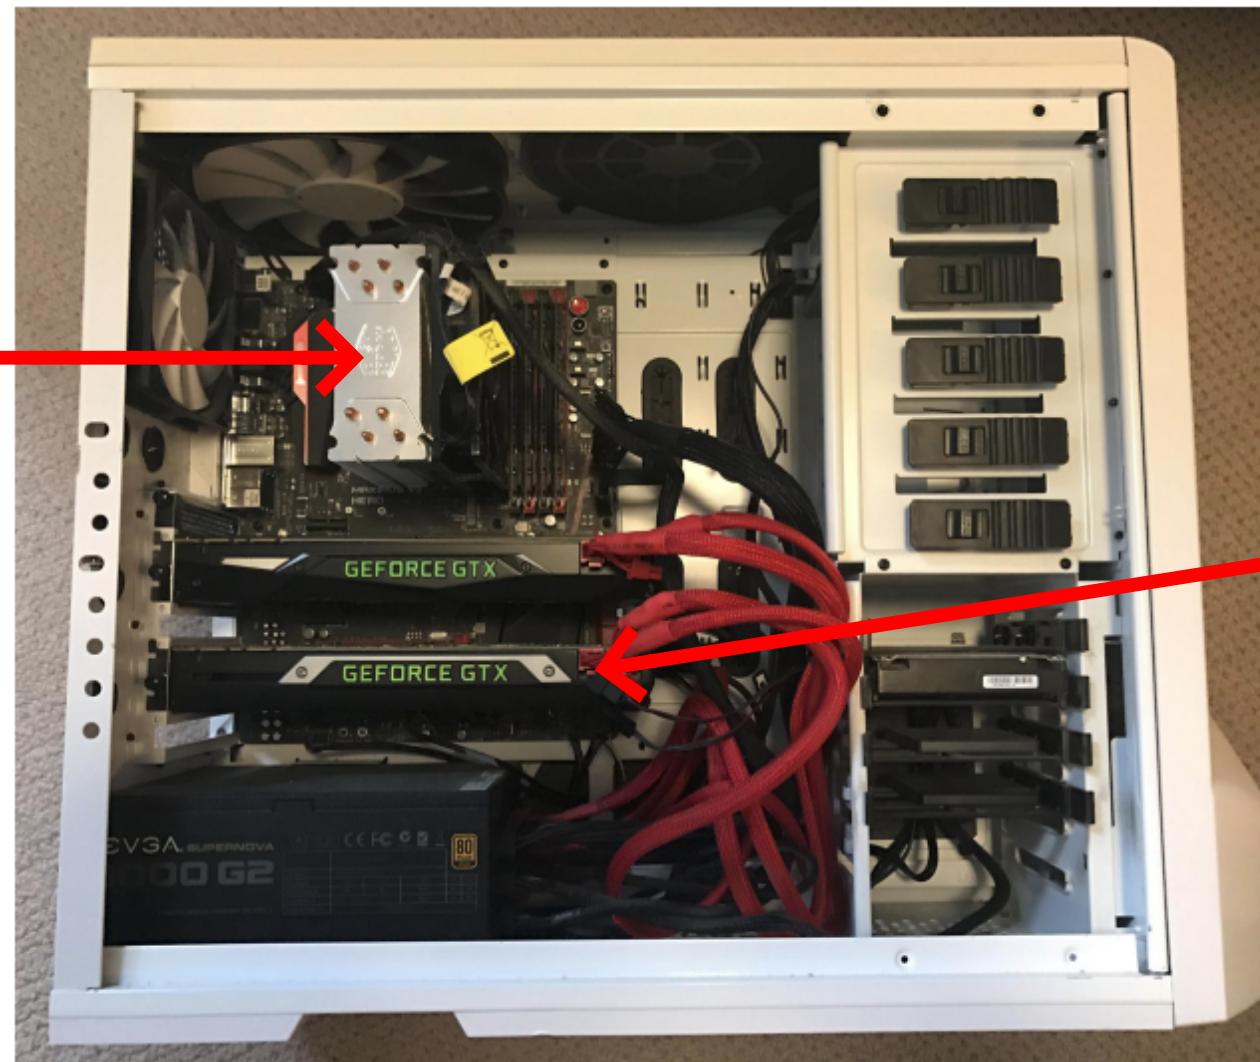
01. CPU vs. GPU

CPU

- Central Processing Unit



This image is licensed under CC-BY 2.0



GPU

- Graphics Processing Unit
- Graphics Card
- takes up a lot of power
- NVIDIA vs. AMD?
- 딥러닝에서는 주로 NVIDIA 사용



01. CPU vs. GPU

CPU vs. GPU in practice

| | # Cores | Clock Speed | Memory | Price |
|-------------------------------------|---|-------------|--------------------|--------|
| CPU (Intel Core i7-7700k) | 4 (8 threads with hyperthreading) | 4.4 GHz | Shared with system | \$339 |
| CPU (Intel Core i7-6950X) | 10 (20 threads with hyperthreading) | 3.5 GHz | Shared with system | \$1723 |
| GPU (NVIDIA Titan Xp) | 3840 | 1.6 GHz | 12 GB GDDR5X | \$1200 |
| GPU (NVIDIA GTX 1070) | 1920 | 1.68 GHz | 8 GB GDDR5 | \$399 |

CPU: Fewer cores, but each core is much faster and much more capable; great at sequential tasks

GPU: More cores, but each core is much slower and “dumber”; great for parallel tasks

CPU

- core 수가 적지만 각각의 core가 독립적으로 작동
- multithreading을 사용하면 최대 20가지의 task 처리 가능
- 대부분의 메모리를 RAM에서 끌어 씀 (RAM: 8, 12, 16, 32GB)

GPU

- core 수가 훨씬 많지만 각각의 clock speed 가 느림
- 독립적x. 많은 core가 하나의 task.
- 행렬곱 연산과 같은 병렬 처리에 유리
- 자체적으로 RAM을 가지고 있음 (RAM-GPU통신 bottleneck 문제)

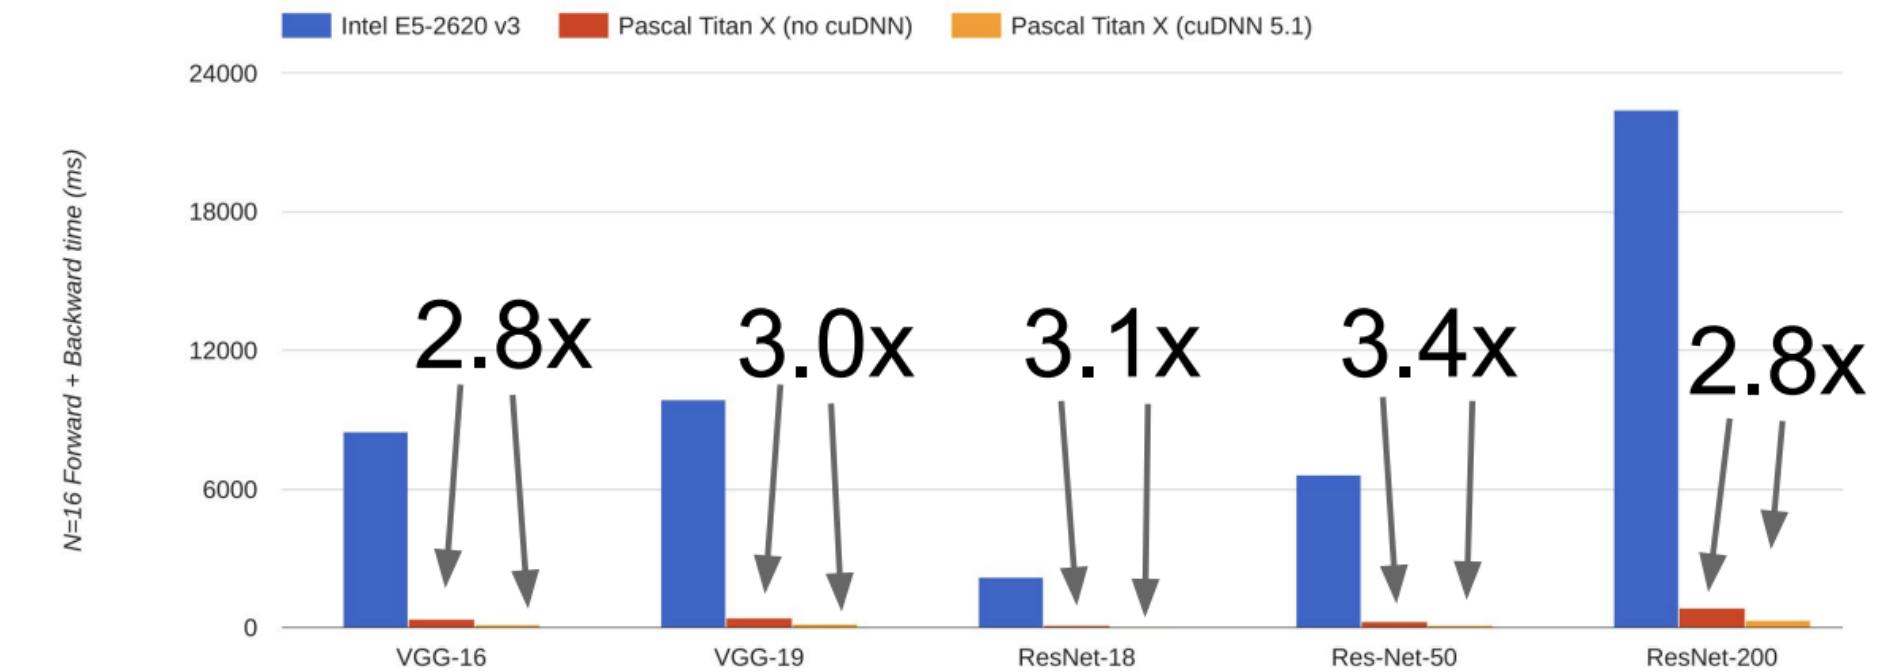
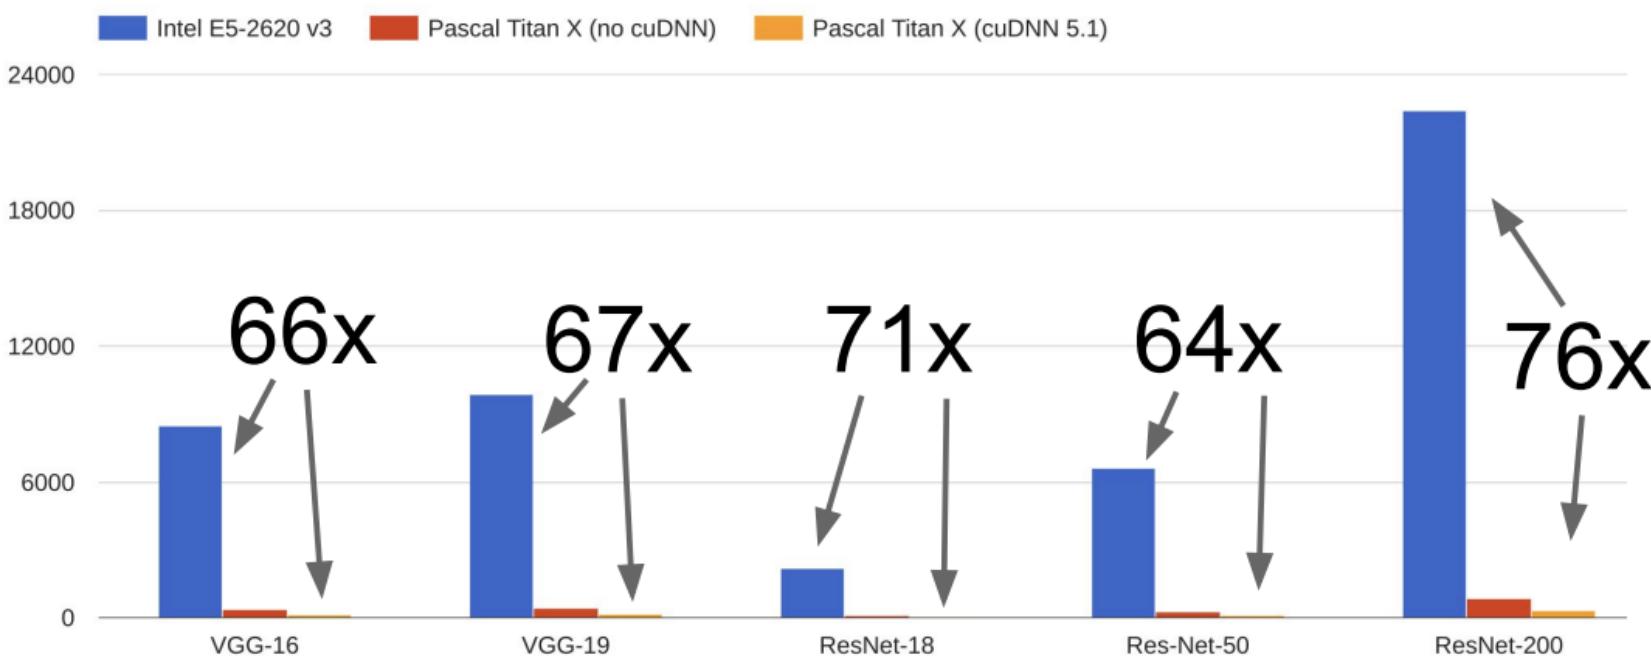
01. CPU vs. GPU

Programming GPUs

- CUDA(NVIDIA only)
 - Write C-like codes that runs directly on the GPU -> 어려움!
 - Higher-level APIs : cuBLAS, cuFFT, cuDNN, etc
 - cuDNN : NVIDIA 하드웨어에서 잘 돌아가는 라이브러리(convolution, forward/backward pass, batch normalization, recurrent networks, etc)를 NVIDIA에서 만들어서 공개한 것.
- OpenCL
 - CUDA와 비슷한데 AMD 하드웨어나 CPU에서도 돌아감.
 - usually slower
- Udacity

01. CPU vs. GPU

CPU vs. GPU in practice



속도에 대한
성능 비교:

GPU > CPU

cuDNN > unoptimized CUDA
(3x performance boost)

02. Deep Learning Frameworks

Today

A bit about these



Mostly these

Paddle
(Baidu)

CNTK
(Microsoft)

MXNet
(Amazon)

Developed by U Washington, CMU, MIT, Hong Kong U, etc but main framework of choice at AWS

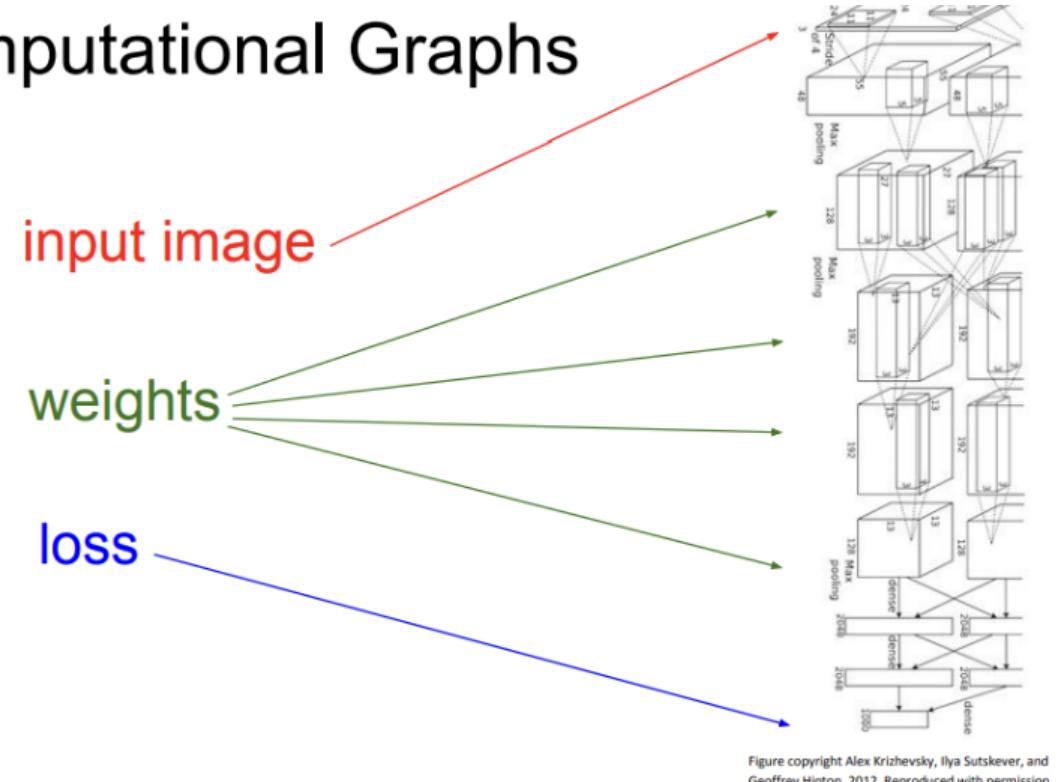
And others...

02. Deep Learning Frameworks

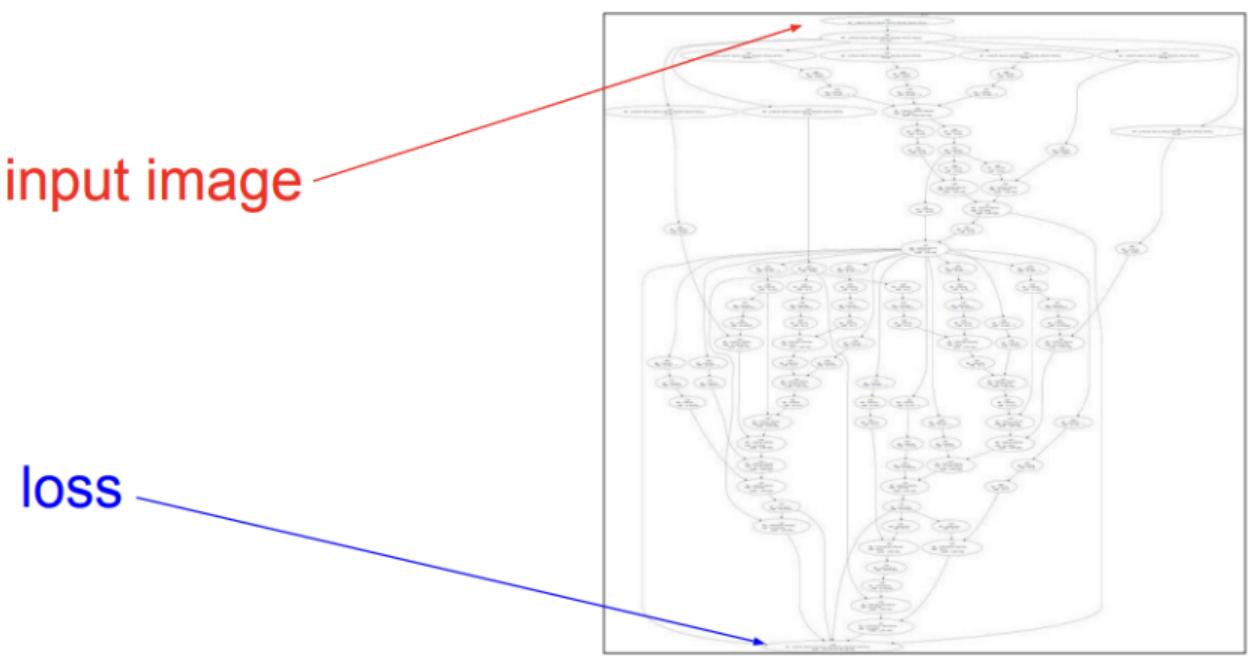
왜 필요한가?

- (1) Easily build computational graphs
- (2) Easily compute gradients in computational graphs
- (3) Run it efficiently on GPU(wrap cuDNN, cuBLAS, etc)

Computational Graphs

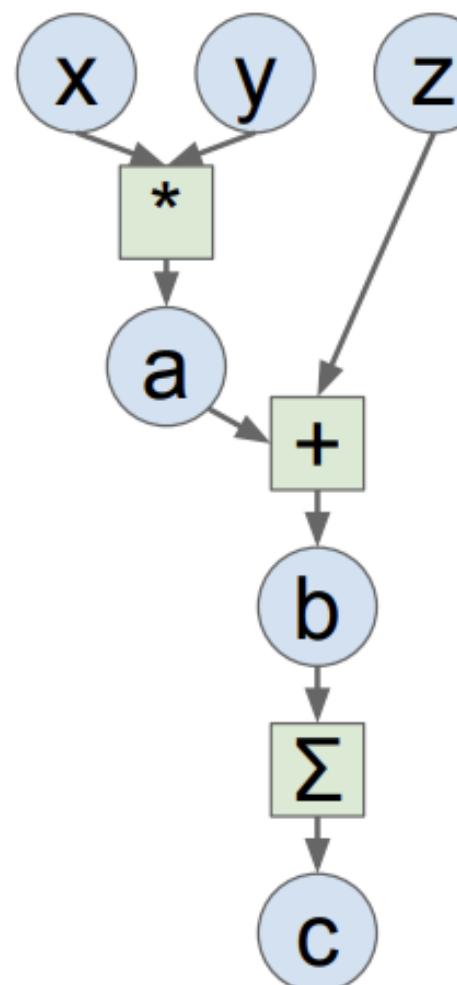


II: Computational Graphs



02. Deep Learning Frameworks

Computational Graphs



Create forward computational graph

TensorFlow

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

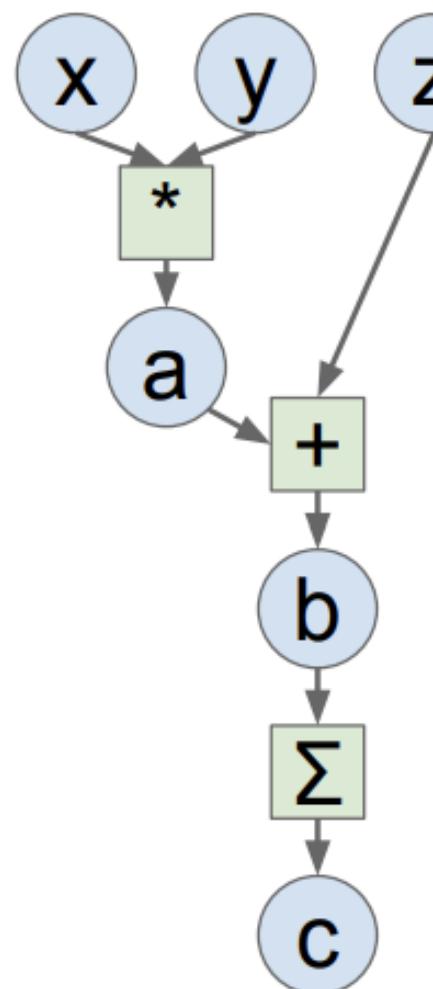
a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

02. Deep Learning Frameworks

Computational Graphs



Ask TensorFlow to
compute gradients

TensorFlow

```
# Basic computational graph
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

x = tf.placeholder(tf.float32)
y = tf.placeholder(tf.float32)
z = tf.placeholder(tf.float32)

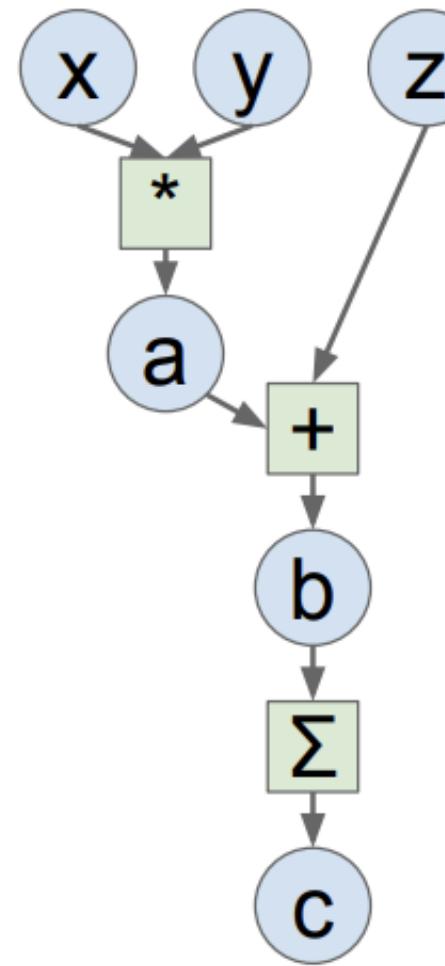
a = x * y
b = a + z
c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

02. Deep Learning Frameworks

Computational Graphs



Tell
TensorFlow
to run on **GPU**

TensorFlow

```
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3000, 4000

with tf.device('/gpu:0'):
    x = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
    z = tf.placeholder(tf.float32)

    a = x * y
    b = a + z
    c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

02. Deep Learning Frameworks

Numpy

```
import numpy as np
np.random.seed(0)

N, D = 3, 4

x = np.random.randn(N, D)
y = np.random.randn(N, D)
z = np.random.randn(N, D)

a = x * y
b = a + z
c = np.sum(b)

grad_c = 1.0
grad_b = grad_c * np.ones((N, D))
grad_a = grad_b.copy()
grad_z = grad_b.copy()
grad_x = grad_a * y
grad_y = grad_a * x
```

TensorFlow

```
import numpy as np
np.random.seed(0)
import tensorflow as tf

N, D = 3, 4

with tf.device('/gpu:0'):
    x = tf.placeholder(tf.float32)
    y = tf.placeholder(tf.float32)
    z = tf.placeholder(tf.float32)

    a = x * y
    b = a + z
    c = tf.reduce_sum(b)

grad_x, grad_y, grad_z = tf.gradients(c, [x, y, z])

with tf.Session() as sess:
    values = {
        x: np.random.randn(N, D),
        y: np.random.randn(N, D),
        z: np.random.randn(N, D),
    }
    out = sess.run([c, grad_x, grad_y, grad_z],
                  feed_dict=values)
    c_val, grad_x_val, grad_y_val, grad_z_val = out
```

PyTorch

```
import torch
from torch.autograd import Variable

N, D = 3, 4

x = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)
y = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)
z = Variable(torch.randn(N, D).cuda(),
             requires_grad=True)

a = x * y
b = a + z
c = torch.sum(b)

c.backward()

print(x.grad.data)
print(y.grad.data)
print(z.grad.data)
```

PyTorch도
거의 비슷함!

numpy와 비슷하고, gradient도 알아서 계산해주고, GPU에서도 돌아간다!

02. Deep Learning Frameworks

Tensorflow

```
import numpy as np  
import tensorflow as tf
```

(Assume imports at the top of each snippet)

TensorFlow: Neural Net

(1) computational graph 정의

First **define**
computational graph

(2) data를 넣고 graph 실행

Then **run** the graph
many times

```
N, D, H = 64, 1000, 100  
x = tf.placeholder(tf.float32, shape=(N, D))  
y = tf.placeholder(tf.float32, shape=(N, D))  
w1 = tf.placeholder(tf.float32, shape=(D, H))  
w2 = tf.placeholder(tf.float32, shape=(H, D))  
  
h = tf.maximum(tf.matmul(x, w1), 0)  
y_pred = tf.matmul(h, w2)  
diff = y_pred - y  
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))  
  
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])  
  
with tf.Session() as sess:  
    values = {x: np.random.randn(N, D),  
              w1: np.random.randn(D, H),  
              w2: np.random.randn(H, D),  
              y: np.random.randn(N, D),}  
    out = sess.run([loss, grad_w1, grad_w2],  
                  feed_dict=values)  
    loss_val, grad_w1_val, grad_w2_val = out
```

02. Deep Learning Frameworks

Tensorflow

TensorFlow: Neural Net

input값(data) 받을 자리 만들기

Create **placeholders** for
input x, weights w1 and
w2, and targets y

x와 w1의 행렬곱 연산 수행(tf.matmul)
+ ReLU 함수 구현(tf.maximum)

loss계산, w1과 w2에 대한 gradient를
tensorflow가 알아서 계산

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))

grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              w1: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad_w1, grad_w2],
                  feed_dict=values)
    loss_val, grad_w1_val, grad_w2_val = out
```

02. Deep Learning Frameworks

Tensorflow

TensorFlow: Neural Net

input값(data) 받을 자리 만들기

Create **placeholders** for
input x, weights w1 and
w2, and targets y

x와 w1의 행렬곱 연산 수행(tf.matmul)
+ ReLU 함수 구현(tf.maximum)

loss계산, w1과 w2에 대한 gradient를
tensorflow가 알아서 계산

just 그래프 "구성".

아직 실행하지 않은 상태!

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))

grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              w1: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad_w1, grad_w2],
                  feed_dict=values)
    loss_val, grad_w1_val, grad_w2_val = out
```

02. Deep Learning Frameworks

Tensorflow

TensorFlow: Neural Net

그래프를 실행시키는 session에 진입

Now done building our graph, so we enter a **session** so we can actually run the graph

그래프에 들어갈 실제 값 생성
(numpy array)

원하는 출력값(numpy array):
loss, grad_w1, grad_w2
feed_dict를 통해 실제 값 전달

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))

grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              w1: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    out = sess.run([loss, grad_w1, grad_w2],
                  feed_dict=values)
    loss_val, grad_w1_val, grad_w2_val = out
```

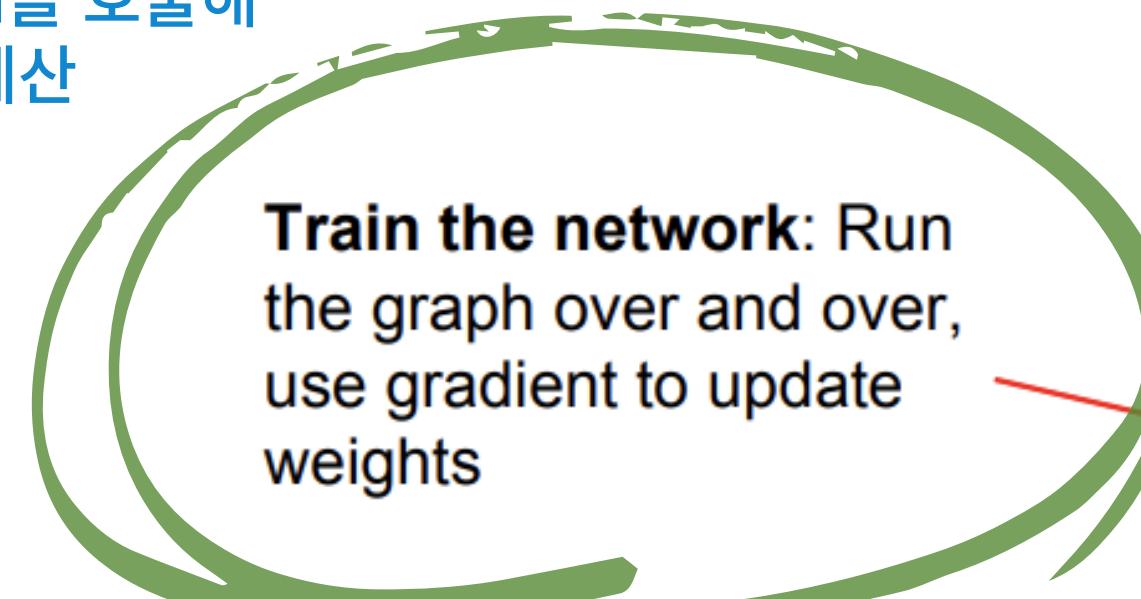
02. Deep Learning Frameworks

Tensorflow

TensorFlow: Neural Net

Problem: copying
weights between CPU /
GPU each step

반복적으로 session.run을 호출해
loss와 grad를 계산



Fei-Fei Li & Justin Johnson & Serena Yeung

Lecture 8 -5151 April 27, 2017

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.placeholder(tf.float32, shape=(D, H))
w2 = tf.placeholder(tf.float32, shape=(H, D))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

with tf.Session() as sess:
    values = {x: np.random.randn(N, D),
              w1: np.random.randn(D, H),
              w2: np.random.randn(H, D),
              y: np.random.randn(N, D),}
    learning_rate = 1e-5
    for t in range(50):
        out = sess.run([loss, grad_w1, grad_w2],
                      feed_dict=values)
        loss_val, grad_w1_val, grad_w2_val = out
        values[w1] -= learning_rate * grad_w1_val
        values[w2] -= learning_rate * grad_w2_val
```

가중치를 업데이트하기 위해 그레디언트를 계산해서
수동으로 Gradient descent에 반영

반복문을 사용하고자 하는 경우

02. Deep Learning Frameworks

Tensorflow

TensorFlow: Optimizer

Can use an **optimizer** to
compute gradients and
update weights

Remember to execute the
output of the optimizer!

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff * diff, axis=1))

optimizer = tf.train.GradientDescentOptimizer(1e-5)
updates = optimizer.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                             feed_dict=values)
```

02. Deep Learning Frameworks

Tensorflow

TensorFlow:
LOSS

Use predefined
common losses

L2 loss를 직접 구하는
코드를 짜지 않아도 됨!

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
loss = tf.losses.mean_squared_error(y_pred, y)

optimizer = tf.train.GradientDescentOptimizer(1e-3)
updates = optimizer.minimize(loss)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                              feed_dict=values)
```

02. Deep Learning Frameworks

Keras: High-level Wrapper

Tensorflow 추상화 패키지

Keras: High-Level Wrapper

딥러닝 아키텍처에서의
대부분의 세부사항을
대신 다뤄줌.

Define model object as
a sequence of layers

optimizer 사용

모델 / loss function specify

코드 한 줄로 train

```
from keras.models import Sequential
from keras.layers.core import Dense, Activation
from keras.optimizers import SGD

N, D, H = 64, 1000, 100

model = Sequential()
model.add(Dense(input_dim=D, output_dim=H))
model.add(Activation('relu'))
model.add(Dense(input_dim=H, output_dim=D))

optimizer = SGD(lr=1e-0)
model.compile(loss='mean_squared_error',
              optimizer=optimizer)

x = np.random.randn(N, D)
y = np.random.randn(N, D)
history = model.fit(x, y, nb_epoch=50,
                     batch_size=N, verbose=0)
```

02. Deep Learning Frameworks

TensorFlow: Other High-Level Wrappers

Keras (<https://keras.io/>)

TFLearn (<http://tflearn.org/>)

TensorLayer (<http://tensorlayer.readthedocs.io/en/latest/>)

tf.layers (https://www.tensorflow.org/api_docs/python/tf/layers)

TF-Slim (<https://github.com/tensorflow/models/tree/master/inception/inception/slim>)

tf.contrib.learn (https://www.tensorflow.org/get_started/tflearn)

Pretty Tensor (<https://github.com/google/prettytensor>)

Sonnet (<https://github.com/deepmind/sonnet>)

Ships with TensorFlow

From Google

From DeepMind

02. Deep Learning Frameworks

Pre-trained models

Examples

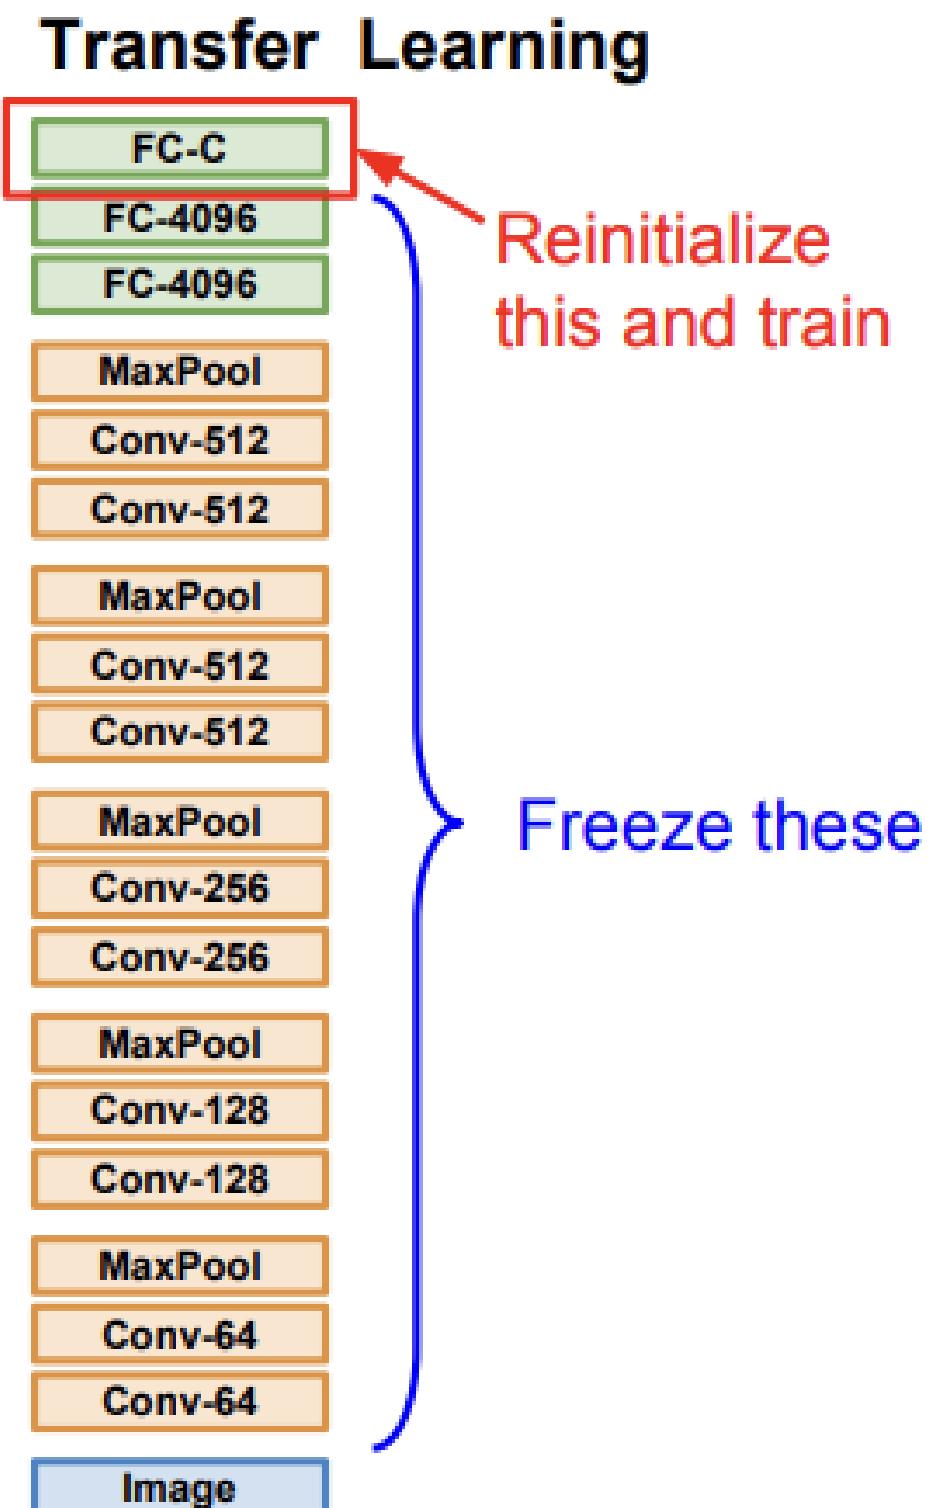
Classify images

```
from resnet50 import ResNet50
from keras.preprocessing import image
from imagenet_utils import preprocess_input, decode_predictions

model = ResNet50(weights='imagenet')

img_path = 'elephant.jpg'
img = image.load_img(img_path, target_size=(224, 224))
x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)

preds = model.predict(x)
print('Predicted:', decode_predictions(preds))
# print: [[u'n02504458', u'African_elephant']]
```



02. Deep Learning Frameworks

Theano

Side Note: Theano

TensorFlow is similar in many ways to **Theano** (earlier framework from Montreal)

```
import theano
import theano.tensor as T

# Batch size, input dim, hidden dim, num classes
N, D, H, C = 64, 1000, 100, 10

x = T.matrix('x')
y = T.vector('y', dtype='int64')
w1 = T.matrix('w1')
w2 = T.matrix('w2')

# Forward pass: Compute scores
a = x.dot(w1)
a_relu = T.nnet.relu(a)
scores = a_relu.dot(w2)

# Forward pass: compute softmax loss
probs = T.nnet.softmax(scores)
loss = T.nnet.categorical_crossentropy(probs, y).mean()

# Backward pass: compute gradients
dw1, dw2 = T.grad(loss, [w1, w2])

f = theano.function(
    inputs=[x, y, w1, w2],
    outputs=[loss, scores, dw1, dw2],
)
```

02. Deep Learning Frameworks

PyTorch의 추상화

PyTorch: Three Levels of Abstraction

GPU에서 작동

Tensor: Imperative ndarray,
but runs on GPU

TensorFlow equivalent

Numpy array

data와 gradient
저장

Variable: Node in a
computational graph; stores
data and gradient

Tensor, Variable, Placeholder

Tensorflow의
higher-level
framework와 유사

Module: A neural network
layer; may store state or
learnable weights

`tf.layers`, or `TFSlim`, or `TFLearn`,
or `Sonnet`, or

02. Deep Learning Frameworks

PyTorch: Tensors

tensor를 사용해 input값 받기

Create random tensors
for data and weights

forward pass

backward pass

gradient와 learning rate을 통해
가중치 계산

```
import torch

dtype = torch.FloatTensor

N, D_in, H, D_out = 64, 1000, 100, 10
x = torch.randn(N, D_in).type(dtype)
y = torch.randn(N, D_out).type(dtype)
w1 = torch.randn(D_in, H).type(dtype)
w2 = torch.randn(H, D_out).type(dtype)
```

```
learning_rate = 1e-6
for t in range(500):
    h = x.mm(w1)
    h_relu = h.clamp(min=0)
    y_pred = h_relu.mm(w2)
    loss = (y_pred - y).pow(2).sum()
```

```
grad_y_pred = 2.0 * (y_pred - y)
grad_w2 = h_relu.t().mm(grad_y_pred)
grad_h_relu = grad_y_pred.mm(w2.t())
grad_h = grad_h_relu.clone()
grad_h[h < 0] = 0
grad_w1 = x.t().mm(grad_h)
```

```
w1 -= learning_rate * grad_w1
w2 -= learning_rate * grad_w2
```

02. Deep Learning Frameworks

자동으로 gradient를
계산하는 함수

but 필요한 대부분의 연산들은
이미 구현되었을 것이기 때문에
autograd을 직접 사용할 일은
잘 없다!

PyTorch: Autograd

We will not want gradients
(of loss) with respect to data

Do want gradients with
respect to weights

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

02. Deep Learning Frameworks

PyTorch: optim

Use an **optimizer** for different update rules

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10

x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out))
loss_fn = torch.nn.MSELoss(size_average=False)

learning_rate = 1e-4
optimizer = torch.optim.Adam(model.parameters(),
                             lr=learning_rate)
for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    optimizer.zero_grad()
    loss.backward()

    optimizer.step()
```

02. Deep Learning Frameworks

PyTorch: nn

Define our model as a sequence of layers

nn also defines common loss functions

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)

model = torch.nn.Sequential(
    torch.nn.Linear(D_in, H),
    torch.nn.ReLU(),
    torch.nn.Linear(H, D_out))

loss_fn = torch.nn.MSELoss(size_average=False)

learning_rate = 1e-4
for t in range(500):
    y_pred = model(x)
    loss = loss_fn(y_pred, y)

    model.zero_grad()
    loss.backward()

    for param in model.parameters():
        param.data -= learning_rate * param.grad.data
```

02. Deep Learning Frameworks

PyTorch: nn Define new Modules

Define our whole model
as a single Module

```
import torch
from torch.autograd import Variable

class TwoLayerNet(torch.nn.Module):
    def __init__(self, D_in, H, D_out):
        super(TwoLayerNet, self).__init__()
        self.linear1 = torch.nn.Linear(D_in, H)
        self.linear2 = torch.nn.Linear(H, D_out)

    def forward(self, x):
        h_relu = self.linear1(x).clamp(min=0)
        y_pred = self.linear2(h_relu)
        return y_pred
```

forward pass

```
N, D_in, H, D_out = 64, 1000, 100, 10

x = Variable(torch.randn(N, D_in))
y = Variable(torch.randn(N, D_out), requires_grad=False)

model = TwoLayerNet(D_in, H, D_out)

criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for t in range(500):
    y_pred = model(x)
    loss = criterion(y_pred, y)

    optimizer.zero_grad()
    loss.backward()
    optimizer.step()
```

loss.backward를 호출하면 매 반복시마다
그래디언트가 자동으로 계산

Fei-Fei Li & Justin Johnson & Serena Yeung

Lecture 8 - $\frac{10}{4}$ April 27, 2017

tensorflow에는 high-level wrapper가 여러 개 소개되었지만 pytorch는
nn뿐!(2017기준...)->그치만 아주 잘 작동함

02. Deep Learning Frameworks

PyTorch: DataLoaders

A **DataLoader** wraps a **Dataset** and provides minibatching, shuffling, multithreading, for you

When you need to load custom data, just write your own Dataset class



```
import torch
from torch.autograd import Variable
from torch.utils.data import TensorDataset, DataLoader

N, D_in, H, D_out = 64, 1000, 100, 10

x = torch.randn(N, D_in)
y = torch.randn(N, D_out)

loader = DataLoader(TensorDataset(x, y), batch_size=8)

model = TwoLayerNet(D_in, H, D_out)

criterion = torch.nn.MSELoss(size_average=False)
optimizer = torch.optim.SGD(model.parameters(), lr=1e-4)
for epoch in range(10):
    for x_batch, y_batch in loader:
        x_var, y_var = Variable(x), Variable(y)
        y_pred = model(x_var)
        loss = criterion(y_pred, y_var)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
```

02. Deep Learning Frameworks

Pretrained models

파이토치
한국 사용자 모임

모델 기여하기
* 현재는 베타입니다 - 앞으로 몇 달 동안 피드백을 수집하고 PyTorch Hub를 개선 예정입니다.

연구자용 –
최신의 연구들을 찾아보고 모델을 확장하세요.

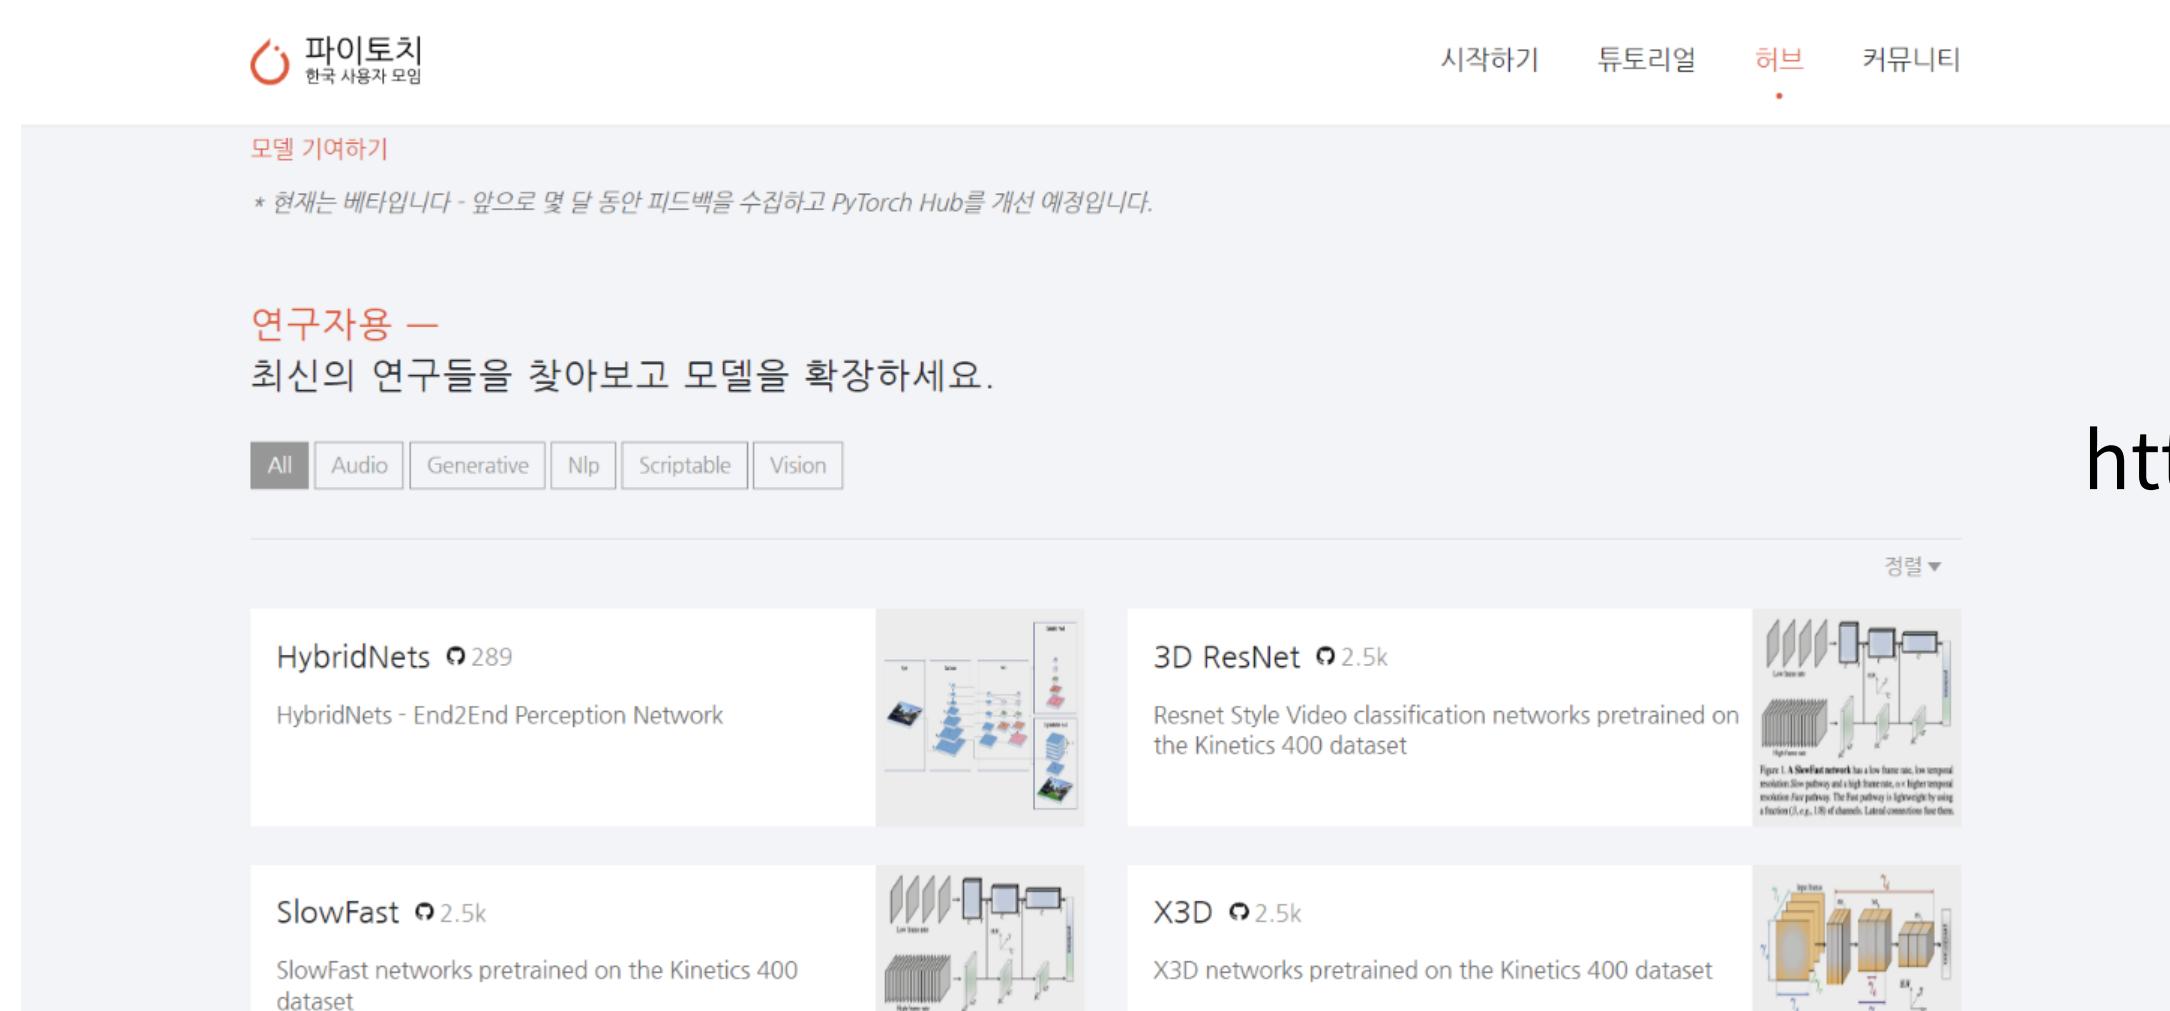
All Audio Generative Nlp Scriptable Vision

HybridNets 289
HybridNets - End2End Perception Network

SlowFast 2.5k
SlowFast networks pretrained on the Kinetics 400 dataset

3D ResNet 2.5k
Resnet Style Video classification networks pretrained on the Kinetics 400 dataset

X3D 2.5k
X3D networks pretrained on the Kinetics 400 dataset



<https://pytorch.kr/hub/>

YOLOv5

By Ultralytics

YOLOv5 in PyTorch > ONNX > CoreML > TFLite

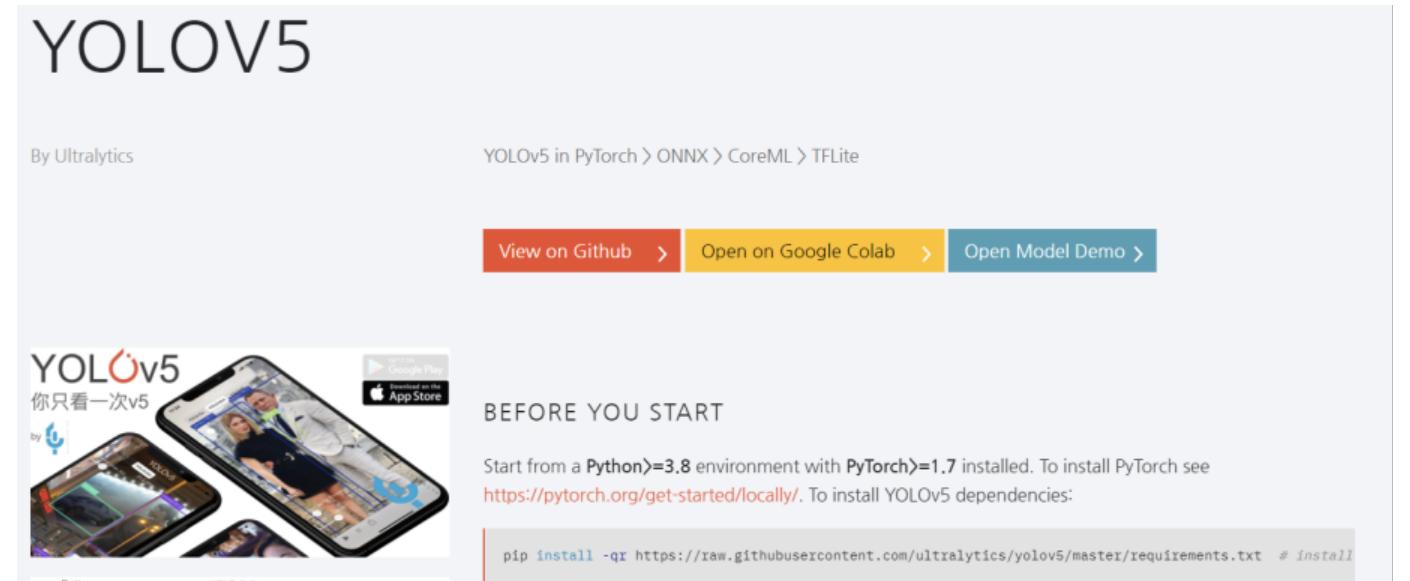
[View on Github](#) [Open on Google Colab](#) [Open Model Demo](#)

YOLov5
你只看一次v5
by Ultralytics

BEFORE YOU START

Start from a **Python>=3.8** environment with **PyTorch>=1.7** installed. To install PyTorch see <https://pytorch.org/get-started/locally/>. To install YOLOv5 dependencies:

```
pip install -qr https://raw.githubusercontent.com/ultralytics/yolov5/master/requirements.txt # install
```



[README.md](#)

Quick Start Examples

▼ Install

Clone repo and install requirements.txt in a **Python>=3.7.0** environment, including **PyTorch>=1.7**.

```
git clone https://github.com/ultralytics/yolov5 # clone
cd yolov5
pip install -r requirements.txt # install
```

▼ Inference

YOLOv5 [PyTorch Hub](#) inference. [Models](#) download automatically from the latest YOLOv5 [release](#).

```
import torch

# Model
model = torch.hub.load('ultralytics/yolov5', 'yolov5s') # or yolov5n - yolov5x6, custom

# Images
img = 'https://ultralytics.com/images/zidane.jpg' # or file, Path, PIL, OpenCV, numpy, list

# Inference
results = model(img)

# Results
results.print() # or .show(), .save(), .crop(), .pandas(), etc.
```

02. Deep Learning Frameworks

Pretrained models

```
pip install torch torchvision torchaudio --extra-index-url https://download.pytorch.org/whl/cu113
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/, https://download.pytorch.org/whl/cu113
Requirement already satisfied: torch in /usr/local/lib/python3.7/dist-packages (1.12.0+cu113)
Requirement already satisfied: torchvision in /usr/local/lib/python3.7/dist-packages (0.13.0+cu113)
Requirement already satisfied: torchaudio in /usr/local/lib/python3.7/dist-packages (0.12.0+cu113)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch) (4.1.1)
Requirement already satisfied: pillow!=8.3.*,>=8.3.0 in /usr/local/lib/python3.7/dist-packages (from torchvision) (7.1.2)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from torchvision) (2.23.0)
Requirement already satisfied: numpy in /usr/local/lib/python3.7/dist-packages (from torchvision) (1.21.6)
Requirement already satisfied: urllib3!=1.25.0,!=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->torchvision) (1.24.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->torchvision) (2022.6.15)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->torchvision) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->torchvision) (2.10)

git clone https://github.com/ultralytics/yolov5
Cloning into 'yolov5'...
remote: Enumerating objects: 12631, done.
remote: Counting objects: 100% (56/56), done.
remote: Compressing objects: 100% (43/43), done.
remote: Total 12631 (delta 26), reused 31 (delta 13), pack-reused 12575
Receiving objects: 100% (12631/12631), 12.34 MIB | 19.68 MIB/s, done.
Resolving deltas: 100% (8681/8681), done.

cd yolov5
pip install -r requirements.txt

install dependencies
import torch #import pytorch
from matplotlib import pyplot as plt #import matplotlib
import numpy as np #import numpy(helps data transformation)
import cv2 #import opencv
```

```
#load model
model = torch.hub.load('ultralytics/yolov5', 'yolov5s')

Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting PyYAML>=5.3.1
  Downloading PyYAML-6.0-cp37-cp37m-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_12_x86_64.manylinux2010_x86_64.whl (596 kB)
Installing collected packages: PyYAML
  Attempting uninstall: PyYAML
    Found existing installation: PyYAML 3.13
    Uninstalling PyYAML-3.13:
      Successfully uninstalled PyYAML-3.13
Successfully installed PyYAML-6.0
```

```
from google.colab import files
uploaded = files.upload()

파일 선택 선택된 파일 없음 Saving friends.jpg to friends.jpg
Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

[ ] img = "friends.jpg"

[ ] results = model(img)
results.print()

image 1/1: 1080x1440 4 persons, 4 cups, 4 forks, 2 knives, 1 dining table
Speed: 89.7ms pre-process, 456.8ms inference, 1.4ms NMS per image at shape (1, 3, 480, 640)

[ ] #rendering
%matplotlib inline
plt.imshow(np.squeeze(results.render())) #imshow 함수로 render
plt.show()
```

```
#rendering
%matplotlib inline
plt.imshow(np.squeeze(results.render())) #imshow 함수로 render
plt.show()
```



02. Deep Learning Frameworks

Torch vs PyTorch

Torch

- (-) Lua
- (-) No autograd
- (+) More stable
- (+) Lots of existing code
- (0) Fast

PyTorch

- (+) Python
- (+) Autograd
- (-) Newer, still changing
- (-) Less existing code
- (0) Fast

Conclusion: Probably use PyTorch for new projects

02. Deep Learning Frameworks

Tensorflow는 computational graph를
한 번 만들어 iterate

Static vs Dynamic Graphs

Tensorflow: static

TensorFlow: Build graph once, then
run many times (**static**)

```
N, D, H = 64, 1000, 100
x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, axis=1))
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

learning_rate = 1e-5
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)
updates = tf.group(new_w1, new_w2)

with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates],
                             feed_dict=values)
```

Build
graph

Run each
iteration

PyTorch: Dynamic

PyTorch: Each forward pass defines
a new graph (**dynamic**)

```
import torch
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10
x = Variable(torch.randn(N, D_in), requires_grad=False)
y = Variable(torch.randn(N, D_out), requires_grad=False)
w1 = Variable(torch.randn(D_in, H), requires_grad=True)
w2 = Variable(torch.randn(H, D_out), requires_grad=True)

learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

New graph each iteration

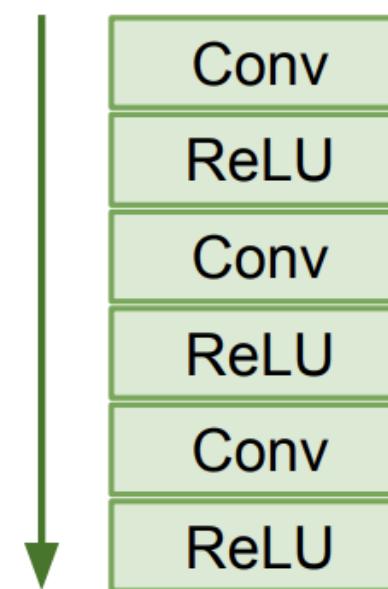
PyTorch는 forward pass 할 때마다 매번
새로운 computational graph를 만듦

02. Deep Learning Frameworks

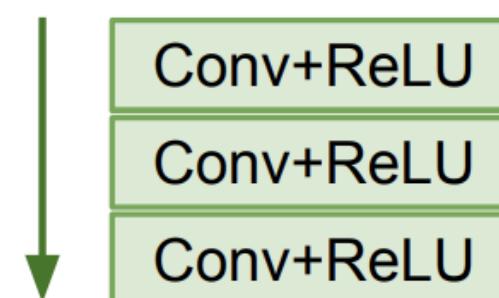
Static vs Dynamic: Optimization

With static graphs, framework can **optimize** the graph for you before it runs!

The graph you wrote



Equivalent graph with **fused operations**



Static vs Dynamic: Serialization

Static

Once graph is built, can **serialize** it and run it without the code that built the graph!

Dynamic

Graph building and execution are intertwined, so always need to keep code around

02. Deep Learning Frameworks

Static vs Dynamic: Conditional

$$y = \begin{cases} w_1 * x & \text{if } z > 0 \\ w_2 * x & \text{otherwise} \end{cases}$$

PyTorch: Normal Python

```
N, D, H = 3, 4, 5  
  
x = Variable(torch.randn(N, D))  
w1 = Variable(torch.randn(D, H))  
w2 = Variable(torch.randn(D, H))  
  
z = 10  
if z > 0:  
    y = x.mm(w1)  
else:  
    y = x.mm(w2)
```

TensorFlow: Special TF control flow operator!

```
N, D, H = 3, 4, 5  
x = tf.placeholder(tf.float32, shape=(N, D))  
z = tf.placeholder(tf.float32, shape=None)  
w1 = tf.placeholder(tf.float32, shape=(D, H))  
w2 = tf.placeholder(tf.float32, shape=(D, H))  
  
def f1(): return tf.matmul(x, w1)  
def f2(): return tf.matmul(x, w2)  
y = tf.cond(tf.less(z, 0), f1, f2)  
  
with tf.Session() as sess:  
    values = {  
        x: np.random.randn(N, D),  
        z: 10,  
        w1: np.random.randn(D, H),  
        w2: np.random.randn(D, H),  
    }  
    y_val = sess.run(y, feed_dict=values)
```

02. Deep Learning Frameworks

Static vs Dynamic: Loops

$$y_t = (y_{t-1} + x_t) * w$$

PyTorch: Normal Python

```
T, D = 3, 4
y0 = Variable(torch.randn(D))
x = Variable(torch.randn(T, D))
w = Variable(torch.randn(D))

y = [y0]
for t in range(T):
    prev_y = y[-1]
    next_y = (prev_y + x[t]) * w
    y.append(next_y)
```

TensorFlow: Special TF control flow

```
T, N, D = 3, 4, 5
x = tf.placeholder(tf.float32, shape=(T, D))
y0 = tf.placeholder(tf.float32, shape=(D,))
w = tf.placeholder(tf.float32, shape=(D,))

def f(prev_y, cur_x):
    return (prev_y + cur_x) * w

y = tf.foldl(f, x, y0)

with tf.Session() as sess:
    values = {
        x: np.random.randn(T, D),
        y0: np.random.randn(D),
        w: np.random.randn(D),
    }
    y_val = sess.run(y, feed_dict=values)
```

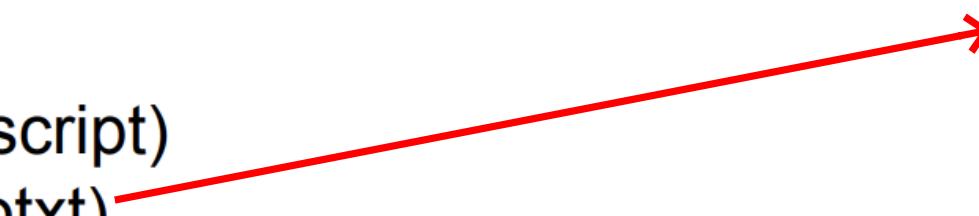
02. Deep Learning Frameworks

Caffe

Caffe: Training / Finetuning

No need to write code!

1. Convert data (run a script)
2. Define net (edit prototxt)
3. Define solver (edit prototxt)
4. Train (with pretrained weights) (run a script)



```
name: "LogisticRegressionNet"
layers {
  top: "data"
  top: "label"
  name: "data"
  type: HDF5_DATA
  hdf5_data_param {
    source: "examples/hdf5_classification/data/train.txt"
    batch_size: 10
  }
  include {
    phase: TRAIN
  }
}
layers {
  bottom: "data"
  top: "fc1"
  name: "fc1"
  type: INNER_PRODUCT
  blobs_lr: 1
  blobs_lr: 2
  weight_decay: 1
  weight_decay: 0
}
inner_product_param {
  num_output: 2
  weight_filler {
    type: "gaussian"
    std: 0.01
  }
  bias_filler {
    type: "constant"
    value: 0
  }
}
layers {
  bottom: "fc1"
  bottom: "label"
  top: "loss"
  name: "loss"
  type: SOFTMAX_LOSS
}
```

02. Deep Learning Frameworks

기업이 사용하는 Framework

Google:
TensorFlow



*“One framework
to rule them all”*

Facebook:
PyTorch + Caffe2



Research



Production

감사합니다!